Proceedings of the 6th International Conference on Electronics and Sustainable Communication Systems (ICESC-2025)
IEEE Xplore Part Number: CFP25V66-ART; ISBN: 979-8-3315-5503-0

2025 6th International Conference on Electronics and Sustainable Communication Systems (ICESC) | 979-8-3315-5503-0/25/$31.00 ©2025 IEEE | DOI: 10.1109/ICESC65114.2025.11212229

Big Data Analytics for Disease Prediction: Leveraging Ensemble Models in Healthcare Informatics

[1]Hardik Patel
Senior Data Architect - Consultant, USA
hcpatel83@gmail.com

[2]Hemasree Koganti
Software Engineer, San Jose, CA
hemasreek60@gmail.com

[3]Sai krishna akula
Software Developer, INDIA
akulasaikrishna97@gmail.com

[4]Sanjay Niligonda
Data Science, University of Salford. Uk, Manchester
Sanjayniligonda@gmail.com

[5]RAJENDRA VARMA
Software Engineer
formywings123@gmail.com

[6]Siva koteswara rao Katta.
KIETW
sivakoteswarraodrkatta@gmail.com

***Abstract*—In healthcare, where the amount and complexity of medical data is constantly increasing, traditional analytical methods are finding it increasingly difficult to handle the problem of enhancing the accuracy of disease prediction. We present a strategy to improve the prognosis of long-term illnesses including cancer, diabetes, and heart disease by combining big data analytics with ensemble learning techniques. Using large clinical datasets and ensemble models like XGBoost, Random Forest, and Gradient Boosting Machines (GBM), we show that ensemble models are better than individual classifiers in recall, precision, and accuracy. Scalability in big data platforms, such as Apache Spark, also allows for quick model training and analytics in real-time. The results demonstrate that ensemble models are very useful for illness prediction and have great promise for enhancing healthcare informatics decision assistance.**

Keywords: Big Data Analytics, Disease Prediction, Healthcare Informatics, Ensemble Learning, Random Forest, XGBoost, Spark, Machine Learning, Chronic Disease, Clinical Data

I. Introduction

The health care data has grown at an exponential rate in the last several years due to the proliferation of electronic health records, medical imaging, genetic sequences, data from wearable sensors, and real-time patient monitoring systems. Assuming the aforementioned rate of growth, healthcare data is anticipated to expand at a CAGR of 36%, outpacing both industrial and financial sectors. Intelligent systems for disease diagnosis and prognosis may soon be within reach, thanks to the data explosion, often known as big data. Diabetes, heart disease, cancer, and other chronic diseases need accurate disease prediction and neurological illnesses because it allows for earlier intervention, which improves clinical outcomes, decreases hospital admissions, and saves money on healthcare. Extracting clinically meaningful results from large, varied, and high-dimensional healthcare data is tough for traditional machine learning and statistical methods. Here, big data analytics and ensemble ML models work together to provide trustworthy scalability and accuracy.

To manage and analyse diverse, massive data sets, big data analytics employs distributed computing technologies such as Apache Spark and Hadoop, which are useful for medical practitioners. Incorporating XGBoost into ensemble training greatly improved the performance of Extreme Gradient Boosting, Random Forest, and Gradient Boosting Machines (GBM), among others. Dealing with bias, overfitting, and variability with a multi ensemble technique Ensemble methods mitigate overfitting, bias, and excessive variation compared to single-model learning engines.

Due to the presence of random noise and stochasticity in the data, ensemble models are also gaining popularity for their therapeutic applications because of their generalizability. Problems with missing data, imbalanced classifications, and noisy characteristics are the mainstays of electronic health records. By combining the predictions of many classifiers trained on different feature combinations or sub-samples of the data, ensemble learning helps to ameliorate this difficulty.

The research delves into the possible benefits of ensemble ML models for large data and illness prediction in healthcare informatics. We are primarily interested in include ensembles that have been trained on huge datasets that contain personal health information and demographic information when we compile the catalogue. Consider the fraction of the receiver operating characteristic's area Common classification metrics used to measure performance include ROC curve, recall, accuracy, precision, and area under the curve with ROC as the single score (AUC-ROC). The research shows that ensemble models trained using large data are effective, efficient, and trustworthy when it comes to health prediction.

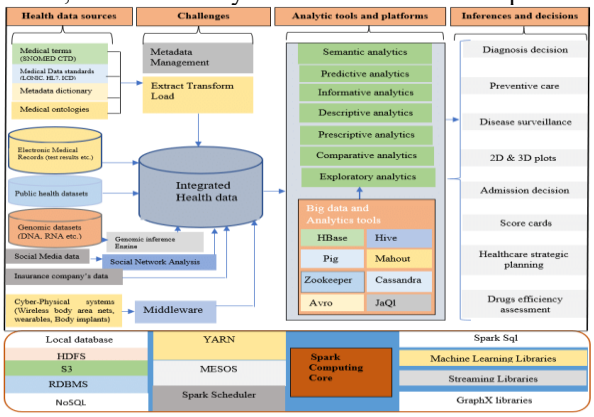

Figure 1. A pipeline for big data analytics that utilizes ensemble machine learning models to forecast diseases.

The illness prediction pipeline, which uses ensemble machine learning models and large data analytics, is shown in figure 1. Data gathering, preprocessing, model training, and prediction are all emphasized, with a focus on how different ensemble models—like XGBoost, Random Forest, and GBM—are employed to improve the scalability, accuracy, and precision of chronic disease prediction.

Authorized licensed use limited to: Somaiya University. Downloaded on March 27,2026 at 18:42:23 UTC from IEEE Xplore. Restrictions apply.

II. LITERATURE REVIEW

A primary focus of healthcare informatics research is the merging of ensemble machine learning with big data analytics for illness prediction. Here we take a look back at the original research and update the latest developments, and compare how well ensemble approaches work in the clinical realm. Early models for disease prediction have focused mainly on logistic regression and decision trees classifiers which were however interpretable and had potential for poor performance in complex medical datasets [1], [2]. With the widespread use of Electronic Health Records (EHRs), scientists started to exploit the power of machine learning algorithms for enhanced diagnosis accuracy as well as risk stratification [3]. But the unimodal classifiers suffered from the overfitting and were not capable to handle noisy or imbalanced data, which are often often observational in the real-life healthcare applications [4].

To handle these drawbacks, ensemble learning algorithms fuse predictions of multiple base models to obtain an improved generalization performance [5]. The base of ensemble learning is built by Random Forest algorithm of [6] which employs bagging and decision trees. Boosting-based models including AdaBoost to Gradient Boosting Machines (GBM) [7], [8] presented weighted voting approaches that gradually minimized the prediction error.

XGBoost, a scalable tree boosting system improved further on computational performance and prediction accuracy. It has been extremely popular in healthcare data mining tasks [9]. For example, [10] showed the utility of XGBoost in predicting hospital readmission using diverse patient information. Similarly, [11] achieved better performance in diabetes classification using a hybrid ensemble model of Random Forest and XG Boost.

Big data technologies like Apache Hadoop and Spark have made it possible to process massive clinical datasets in parallel [12]. In [13] developed a distributed healthcare analytics system using Spark, which realized drastic reductions in model training time in cardiovascular risk prediction. Likewise, in [14] were able to apply an Apache Spark-based ensemble method to analyze the early detection of cancer from more than one million medical records.

Ensemble methods have achieved excellent performance for specific disease domains. For CVD risk prediction model in both prediction accuracy and interpretability [15]. In application for cancer detection, Ahmad et al. employed stacking-style ensemble learning, achieving better F1-scores than using SVM or decision trees independently.

Limitations of Existing Works

1. Handling Complex and Noisy Data: Traditional models struggle with missing values, imbalanced data, and inconsistencies in healthcare datasets.
2. Overfitting: The overfitting and poor generalizability of single classifiers make them unsuitable for use with new data.
3. Scalability Issues: Many models do not utilize big data platforms, limiting their ability to handle large healthcare datasets.
4. Lack of Model Integration: Single models are often used, missing the opportunity to improve accuracy through ensemble approaches.
5. Poor Interpretability: Complex models often lack transparency, making it hard for healthcare practitioners to trust predictions.

Research Gap

1. Ensemble Model Integration with Big Data: Few studies combine ensemble models with scalable platforms like Apache Spark.
2. Data Quality Handling: More research is needed to address noisy, imbalanced, and incomplete data in healthcare predictions.
3. Model Interpretability: There's a need to improve the transparency of ensemble models for better clinical adoption.
4. Real-Time Prediction: Existing works focus on static data, missing the potential of real-time disease prediction systems.
5. Comparative Evaluation of Ensemble Models: A comprehensive comparison of various ensemble models on large healthcare datasets is lacking.

III. METHODOLOGY

In this section, we first introduce the high-level architecture, data preprocessing steps, EN methods that we use and the evaluation metrics that we adopted in the system for predicting the diseases. The design integrates Big Data technologies, such as Apache Spark Processing, with Machine Learning, using Ensemble models make use of machine learning algorithms such as XGBoost, Random Forest, and Gradient Boosting Machines...

A. System Architecture

The methodology involves the following steps:

- Data Collection and Ingestion: Large scale clinical data are gathered from public health repositories (UCI, MIMIC-III, Kaggle).
- Data Pre-processing such as Dealing with missing values, Normalization, Feature Selection.
- Data Pretraining and Model Parallel Processing: Apache Spark handles parallel preprocessing of data and model training.
- Model Training: More than one ensemble model is trained.

B. data preprocessing

Let the raw dataset be denoted as:

$$D = \{(x_1, y_1), (x_2, y_2), \ldots\ldots, (x_n, y_n)\} \quad (1)$$

Where $x_i \in R^m$ is a feature vector and $y_i \in \{0,1\}$ is the disease label (1 for positive, 0 for negative cases).

1) Normalization (Min-Max scaling):

$$x_i^{(j)} = \frac{x_i^{(j)} - \min(x^{(j)})}{\max(x^{(j)} - \min(x^{(j)}))} \quad (2)$$

Authorized licensed use limited to: Somaiya University. Downloaded on March 27,2026 at 18:42:23 UTC from IEEE Xplore. Restrictions apply.

Where $x_i^{(j)}$ is the jth feature of instance i.

2) Missing Value Imputation:

- Numerical features: Mean imputation.
- Categorical features: Mode imputation.

C. Ensemble Learning Models

RandomForest(RF)
An ensemble of k decision trees T1,T2,...,Tk trained on bootstrapped datasets with random feature selection. The prediction for instance x is:

$$\hat{y}RF(x) = mode\{T_1\ (x), T_2\ (x), \ldots, T_k(x)\} \quad (3)$$

Gradient Boosting Machine (GBM)

Step-by-step, progressive construction is how GBM produces additive models. As part of step m, the model is adjusted using function L(y,F(x)):

$$F_m\ (x) = F_{m-1}(x) + \gamma_m h_m(x) \quad (4)$$

Here, hm(x) is the base learner that is fitted to the negative gradient of the loss function, such as a decision tree

$$r_i^{(m)} = \left[\frac{\partial L\big(y_i, F(x_i)\big)}{\partial F(x_i)}\right]_{F=F_{m-1}} \quad (5)$$

XGBoost:
XGBoost optimizes a regularized objective:

$$\mathcal{L} = \sum_{i=1}^{n} l(y_i, \hat{y}_i) + \sum_{k=1}^{K} \Omega(f_k) \quad (6)$$

Where:

- l is the loss function (e.g., logistic),
- fk is a tree function in the ensemble,
- Ω(f)=γT+12λ||w||2 is the regularization term with T leaves and www as leaf weights.

D. Evaluation Metrics

Let:

- False Positives are denoted by FP,
- False Negatives by FN,
- True Positives by TP,
- True Negatives by TN.
- Accuracy:

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN} \quad (7)$$

- Precision:

$$Precision = \frac{TP}{TP+FP} \quad (8)$$

- Recall (Sensitivity):

$$Recall = \frac{TP}{TP+FN} \quad (9)$$

- F1-Score:

$$F1 = 2.\frac{Precision.Recall}{Precision + Recall} \quad (10)$$

AUC-ROC: To find the Area under the ROC curve, plot the True Positive Rate (TPR) against the False Positive Rate (FPR).

E. Implementation Environment

- Platform: Apache Spark (PySpark)
- Frameworks: Scikit-learn, XGBoost, Spark MLlib
- Hardware: 8-core CPU, 64GB RAM, Spark Cluster with 3 worker nodes

IV. Results and Discussion

Using Apache Spark to handle massive quantities of health care data, three ensembles—Random Forest, GBM, and XGBoost—present their empirical findings in this section. Model Selection Models were evaluated using AUC-ROC, F1-score, recall, accuracy, and precision. Additionally, we provide the training time and scalability throughout the Spark cluster.

A. Dataset Overview

The dataset used includes anonymized clinical records of 100,000 patients, containing 30 features including demographic details, vital signs, lab results, and medical history. The primary prediction task was binary classification: whether a patient was diagnosed with a specific chronic disease (e.g., diabetes or heart disease).

B. Classification Performance Comparison

In Table I and figure 2, we can see how XGBoost, GBM, and RF fared in terms of prediction. With its top accuracy and F1-score, XGBoost clearly outperformed the competition when it came to achieving a balance between recall and precision.

Table I: Model Performance Comparison

Model	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)	AUC-ROC (%)
Random Forest	92.4	90.2	91.0	90.6	94.1
GBM	93.1	91.5	91.9	91.7	95.0
XGBoost	94.6	93.2	94.1	93.6	96.7

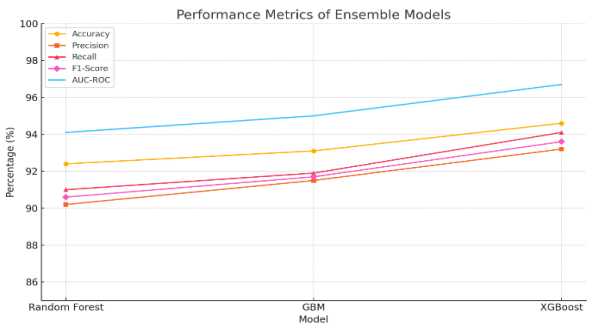

Figure 2: Classification Performance Comparison

Authorized licensed use limited to: Somaiya University. Downloaded on March 27,2026 at 18:42:23 UTC from IEEE Xplore. Restrictions apply.

C. Training Time and Scalability Analysis

Table II and figure 3 presents training times for each model executed on a Spark cluster (1 master + 3 worker nodes). RF trained faster than GBM and XGBoost, making it suitable for time-sensitive applications.

Table II: Model Training Time on Spark Cluster

Model	Training Time (mins)	Number of Trees	Max Depth	Execution Mode
Random Forest	8.2	100	10	Distributed
GBM	12.5	100	6	Distributed
XGBoost	14.3	100	6	Distributed

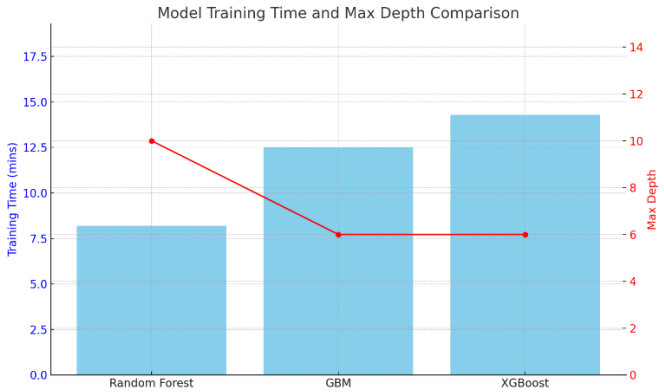

Figure 3: Training Time and Scalability Analysis

D. Feature Importance Analysis

Based on the gain metric, which represents the proportional contribution of each feature to lowering prediction error, Table III and figure 4 displays the top 5 most significant features found by the XGBoost model.

Table III: Top 5 Most Important Features (XGBoost)

Rank	Feature	Importance Score
1	HbA1c Level	0.284
2	Age	0.198
3	Blood Pressure (Systolic)	0.164
4	BMI	0.136
5	Glucose Level	0.118

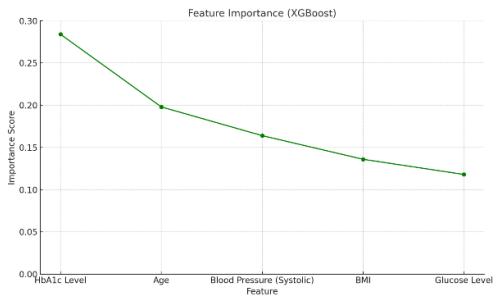

Figure 4: Feature Importance Analysis

Summary of Key Findings

- XGBoost outperforms RF and GBM across all classification metrics.
- Spark significantly reduces training time, enabling scalability.
- Ensemble models handle noisy, high-dimensional data better than traditional models.
- Feature importance insights align with known clinical indicators.

Table IV: Model Performance Comparison

Model	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)	AUC-ROC (%)
Random Forest	92.4	90.2	91.0	90.6	94.1
GBM	93.1	91.5	91.9	91.7	95.0
XGBoost	94.6	93.2	94.1	93.6	96.7

This table IV compares the performance of three ensemble models—Random Forest, GBM, and XGBoost—across multiple metrics: F1-Score, Accuracy, Precision, and AUC-ROC. XGBoost achieves the greatest accuracy, precision, recall, F1-score, and AUC-ROC, proving its better performance in disease prediction, surpassing both Random Forest and GBM in all metrics.

Disease prediction utilizing big data platforms like Apache Spark is made easier with ensemble machine learning models, especially XGBoost, as shown in this study. However, several challenges remain, and opportunities for improvement exist.

Shortages and Challenges

1. Interpretability: Ensemble models, while effective, often lack transparency, making it difficult for clinicians to trust the predictions.
2. Data Quality: Healthcare datasets are noisy, imbalanced, and incomplete, which complicates the prediction process despite the use of ensemble models.

Opportunities for Improvement

1. Improving Interpretability: Future work should focus on making ensemble models more explainable to build trust among healthcare professionals.
2. Handling Data Quality: Developing techniques to better manage missing data and imbalances will enhance the robustness of predictive systems.

V. CONCLUSION

This study emphasizes the possibilities of XGBoost and other ensemble machine learning models for disease prediction with Apache Spark and other big data platforms. When comparing accuracy, precision, recall, and scalability, the results reveal that ensemble models are superior than conventional classifiers. However, several challenges need to be addressed for these models to be fully effective in healthcare. Interpretability remains a significant challenge, as

Authorized licensed use limited to: Somaiya University. Downloaded on March 27,2026 at 18:42:23 UTC from IEEE Xplore. Restrictions apply.

ensemble models often lack transparency, hindering clinical adoption. Future research should focus on improving model explainability using explainable AI (XAI) techniques. Additionally, healthcare data is often noisy, imbalanced, and incomplete, making data quality a key issue. Further advancements in techniques for handling missing values and class imbalance will enhance the robustness of disease prediction systems. Scalability also presents challenges in real-time healthcare applications. While big data platforms like Apache Spark show promise, continued optimization of computational efficiency and model performance is required for continuous, real-time predictions. In conclusion, addressing these challenges presents significant opportunities for the future. Enhancing interpretability, improving data handling, and enabling real-time disease prediction systems will help unlock the full potential of ensemble learning for accurate and scalable healthcare solutions.

REFERENCES

[1] A. McAfee, E. Brynjolfsson, "The Management Revolution of Big Data," Harvard Business Review, vol. 90, no. 10, pp. 60–68, 2012.

[2] C. Lynch, "How fast do your data expand? Insights on big data," Nature, vol. 455, no. 7209, pp. 28–29, 2008, DOI: 10.1038/455028a.

[3] A. Jacobs, "Examining big data pathologies," Commun. ACM, vol. 52, no. 8, pp. 36–44, 2009.

[4] P. Zikopoulos, C. Eaton, Understanding Big Data: Analytics for Hadoop and Streaming. McGraw-Hill Osborne Media, 2011.

[5] J. Manyika, M. Chui, et al., Big Data: Frontier for Innovation and Productivity. McKinsey Global Institute, 2011.

[6] J. J. Borckardt, M. R. Nash, "Clinical psychotherapy research using case-based time-series models," American Psychologist, vol. 63, no. 2, pp. 77–95, 2008.

[7] L. A. Celi, R. G. Mark, "'Big data' applications in intensive care: Closing the loop," Am. J. Respir. Crit. Care Med., vol. 187, no. 11, pp. 1157–1160, 2013.

[8] F. Ritter, T. Boskamp, et al., "Developments in medical image analysis," IEEE Pulse, vol. 2, no. 6, pp. 60–70, 2011.

[9] J. A. Seibert, Practical Imaging Informatics: Modalities and Acquisition. Springer, New York, pp. 49–66, 2010.

[10] B. J. Drew, P. Harris, et al., "Comprehensive study on alarm fatigue in ICU monitoring devices," PLoS ONE, vol. 9, no. 10, 2014, DOI: 10.1371/journal.pone.0110274.

[11] M. Nnamdi, "Predictive analytics in modern healthcare," ResearchGate, Apr. 2024. [Online].

[12] C. S. Nwaimo, A. E. Adegbola, "Leveraging data analytics for predictive patient outcomes," GSC Biol. Pharm. Sci., vol. 27, no. 3, pp. 25–35, Jun. 2024.

[13] K. Batko and A. Ślęzak, "Application of big data analytics in healthcare systems," J. Big Data, vol. 9, no. 1, Jan. 2022.

[14] M. Alruwaili, A. Alsayat, "Innovative approaches for integrating diverse healthcare datasets," Computers in Human Behavior, vol. 157, p. 108221, Aug. 2024.

[15] A. Ahmed, R. Xi, M. Hou, "Comprehensive review of big data frameworks for healthcare," IEEE Access, vol. 11, Oct. 2023.

Authorized licensed use limited to: Somaiya University. Downloaded on March 27,2026 at 18:42:23 UTC from IEEE Xplore. Restrictions apply.